

Automating Code Magnet Generation

Julia Dana

Contents

1	Introduction	4
1.1	The Overview: Easier Creation of Code Magnet Microlabs	4
1.2	The Context: What is WAGS?	4
1.3	The Problem: Brittle Input	4
1.4	The Solution: Parsing by Grammar	5
2	Development results and future extensions	7
2.1	What It Does: Internal Functions	7
2.1.1	Parses by Grammar	7
2.1.2	Alternative Magnets and Other Instructor Directives . . .	8
2.1.3	Improved Representation of Magnets	9
2.2	What It Does: External Functions	10
2.2.1	Command Line Tool	10
2.2.2	Automated Interaction With WAGS Website	10
2.3	What It Could Do: Future Extensions	10
2.3.1	Automatic Generation of Alternative/Distractor Magnets	10
2.3.2	GUI Tool	10
2.3.3	Further Automated Interaction With WAGS Website . .	11
2.3.4	Supporting New Languages	11
2.3.5	Additional Serialization Formats	11
3	User Guide	11
3.1	Quick Start - CLI Tool	11
3.2	Installation for Use	12
3.3	Explanation of Magnet Creation	12
3.3.1	Java	12
3.3.2	Python 3	12
3.4	Modifying the Magnet Creation: Instructor Directives	12
3.4.1	ALT and ENDALT: Creating Alternative Magnets	13
3.4.2	NODROP: Suppressing Drop Zones	13
3.4.3	EXTRAMAG: Creating an Extra Magnet	13
4	Developer Notes	13
4.1	The Environment: Languages and Libraries Used	13
4.1.1	Installation for Development	13
4.2	The Design	14
4.2.1	Directory Structure	14
4.2.2	JRuby / Java Integration	14
4.2.3	Key Classes	14
4.3	The Testing	14
4.4	The Configuration: Specifying Magnet Sections	14
4.5	The Expansion: Adding a New Language	15
4.5.1	New Grammar Specification	15
4.5.2	Rakefile: Useful Actions for Adding a New Language . . .	16

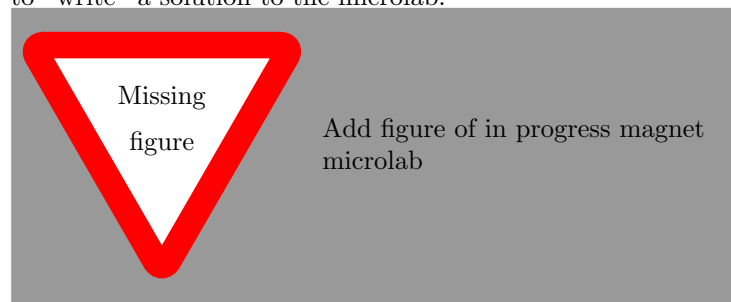
1 Introduction

1.1 The Overview: Easier Creation of Code Magnet Microlabs

The purpose of this project is to assist in the creation of code magnet microlab assignments for WAGS by creating magnets from a completed solution file. This is accomplished in a manner that supports multiple programming languages, and allows additional languages to be added with minimal configuration. Additional tools that support this idea of easier creation of assignments are also included, such as an automated interaction with the WAGS website to create assignments. Also, this project defines new formats for representing magnets, both as objects, and serialized to JSON or YAML.

1.2 The Context: What is WAGS?

WAGS (Web Automated Grading System) is an ongoing project of Appalachian State University. It is an online tool for microlabs. Microlabs are short, 5-10 minutes hands-on activities that are intended to be done as a part of a regular (i.e. not lab) class session to reinforce the concepts that are currently being covered. There are multiple types of microlabs provided by WAGS, but the one that this project is interested in is code magnet microlabs. These are microlabs where the student is given code magnets (pieces of code). Then the student must choose the correct magnets to use and drag and drop them into the correct order to “write” a solution to the microlab.



1.3 The Problem: Brittle Input

The problem is that creating these magnet microlab assignments on the WAGS website is a somewhat painful process. The current parser creates magnets based on whitespace (line breaks and indentation). This limits the style of code that can be used. For example, opening curly braces in Java must be on the same line as the class/method/loop they are associated with. Also, statements have to be on a single line, regardless of length. See Figure 1 for an example [4]. These stylistic limitations, while limiting, do not force your input to the microlab creation to become an invalid solution file. However, to create alternative magnets for the existing parser, one adds extra lines to the input

The iterativeFibonacci in Standard Format

```
public int iterativeFibonacci(int num){  
    if(num == 0 || num == 1) {  
        return num;  
    }  
    int current = 1;  
    int previous = 0;  
    int temp;  
    while (num > 1){  
        temp = current;  
        current = current + previous;  
        previous = temp;  
        num--;  
    }  
    return current;  
}
```

One statement per line

Put the opening brace on the same line as the statement.


Put the closing brace on a separate line.

You can add comments using // at the end of the lines.

Figure 1: Putting an Input File In Standard Format

file. This does change the input file to no longer be a correct solution, and can even change it to no longer be a valid file for its language. This means that the instructor now has to maintain and keep in sync both a proper solution and the WAGS input file. See Figure 2 to see how adding alternative magnets can change an input file[4].

If your input cannot be handled by this brittle parser. WAGS does provide a manual input for magnets. However, using this manual input requires magnets to either be entered one at a time to the magnet creation wizard, or for the user to directly type the final magnet (including HTML escape sequences) to the parsed output section.



1.4 The Solution: Parsing by Grammar

The solution to the problems raised by the old, brittle parser is to use a more robust parser that is based on the grammar of the language, rather than the

Adding Alternative Magnets

```

public int iterativeFibonacci(int num){
    if(num == 0 || num == 1) {
        return num;
        return 1;
    }
    int current = 1;
    int previous = 0;
    int previous = 1;
    int temp;
    while (num > 1){
    while (num >= 1){
        temp = current;
        current = current + previous;
        previous = temp;
        previous = current;
        num--;
        num++;
    }
    return current;
    return temp;
}

```

First put your method in standard format.

Duplicate a line you want an alternative magnet and make the desired change. Alternative magnets are shown in red for the Fibonacci method.

The order of the lines of code is not important; one magnet is made for each line. But for readability by another human, it is recommended to put the correct magnets in order and to put alternative magnets next to the correct magnet.

Figure 2: Adding Alternative Magnets

formatting of the input file. This project does this by using the ANTLR4 parser generator[7] and freely available grammars for common languages[2].

A grammar is a formal specification of a language. It is a series of rules (called productions) that define valid phrases of a language. A single production could define the structure of an entire file, or just be a single statement or expression. A production is the valid components of a statement (tokens and other rules). A rule can match one or more productions.

There are many parallels between a grammar of a programming language, and the grammar of a natural language (such as English). As a rough comparison, using a grammar-based parser to break code into magnets is like breaking apart a paragraph based on sentences or phrases, rather than on every period (Which commonly indicates the end of a sentence, but also has other uses, such as abbreviations). This type of parser will create tokens from the input, which are loosely the equivalent of words, and then break them into (sometimes nested) phrases and sentences according to rules called productions. Rules are named, and will sometimes contain more than one production as alternatives that provide the same function in a language.

A parser uses these rules to build a parse tree. The root of the parse tree is the start rule. The interior nodes of the parse tree are rules from the grammar. The leaf nodes of the parse tree are tokens from the lexer.

This project specifies which rules in the grammar should create their own magnets. If a rule is nested in another rule only the magnet of the nested rule gets the text. The magnet for the outer rule replaces the text of the nested rule with a drop zone (an area that accepts magnets)

Using a parsed file gives us a high level of control over how and when to

create magnets. However, writing a robust parser is a non-trivial task. Today, most parsers are created by parser generators. These take the description of the language given in a grammar, and creates the parser for that language. There are many such parser generators, including YACC [6], Bison [3], and CUP [5]. This project uses the ANTLR4 [1] parser generator, which generates parsers in Java.

expand on
ANTLR -
LL(star) -
other fea-
tures

2 Development results and future extensions

2.1 What It Does: Internal Functions

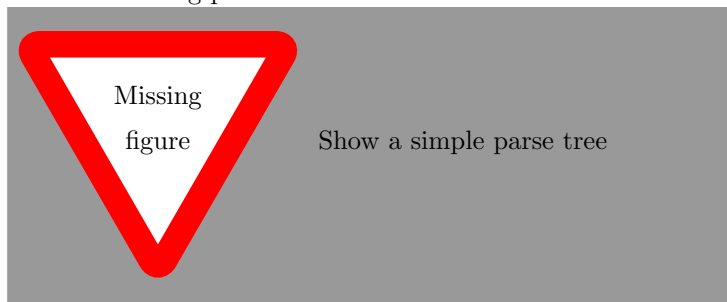
2.1.1 Parses by Grammar

The magnetizer takes an input file, and parses it with the parser generated by ANTLR from the grammar specification. The result of this parsing is a parse tree, which is then used to create magnets. ANTLR can show us a graphical representation of this tree, as follows.

For this simple Java file, Hello.java



The resulting parse tree is:



ANTLR provides two mechanisms for traversing this tree once it is created. One is a walker, which triggers enter and exit methods in a listener for each type of internal node, and the other is a visitor, which has methods for visiting each type of internal node. For both mechanisms, ANTLR provides a Java class which is a base version of the listener or visitor, with stubs of the methods containing minimal functionality. The programmer can then subclass these base versions


to perform their project-specific actions at certain nodes. This project initially used the listener mechanism, but later switched to using the visitor to make it easier to implement the NODROP instructor directive. One key difference between the mechanisms is that the walker/listener method will always visit every node in the tree. With the visitor method, the base (ANTLR-generated) visitor visits every node by default, but the specific visitor can choose to not visit some or all of the children of a node, because the visit method is responsible for choosing when to visit children.

The visitors for this project are dynamically created JRuby subclasses of the ANTLR-generated base visitors. There is an ANTLR-generated base visitor per supported language. There is a single JRuby generator that creates a subclass of the Java base visitor for any language. The generation of the subclass is controlled by an entry in a configuration file which specifies the name of the grammar used (language name), and which rules (which become internal nodes in the parse tree) trigger the creation of a magnet. This dynamic generation prevents repeated code because we want to perform the same types of actions at every node that triggers a magnet, but the structure that ANTLR expects requires separate methods.

One difficulty encountered in creating magnets from the parse tree is that the ANTLR-generated parser (like most parsers) discards whitespace once the parse tree has been built. By default, only the tokens themselves are kept. So, although the nodes in the tree provide a `getText()` method, the result of this is something like `publicclassMyClass` rather than the expected `public class MyClass`. However, ANTLR provides a way to keep extra information such as whitespace and comments around for later use if they are needed called hidden channels 2.1.1. This splits the incoming token stream into a channel for use by the parser, and hidden channels that can later be accessed to recreate the original text (with formatting) for the magnet. This hidden channel mechanism is also used to implement instructor directives.

Show Java configuration file

Show subclass generating code



2.1.2 Alternative Magnets and Other Instructor Directives

Instructor directives allow instructors to control how magnets are created. They are implemented as special comments that go to a special token channel. The

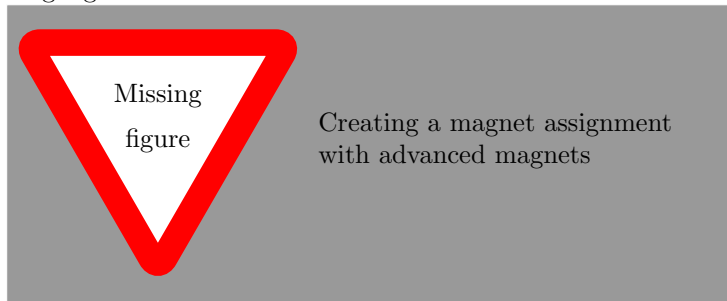
content of directives are the same across all languages, but the triggering comment syntax is specific to each input language. Current directives implemented all instructors to suppress drop zones (NODROP), to indicate that a duplicate magnet with a section of alternate text should be created (ALT and ENDALT), and to create an unrelated magnet (EXTRAMAG).

All the directives currently work with Java input file. However, there are some limitations with directives in Python. Because comments in Python always go to the end of a line, creating a magnet with alternate text does not work, since the original text cannot be surrounded by comments. Also because of line-based comments, unrelated magnets must be created in a single line. There are also some cases where directives do not properly get connected to the expected node in the tree because indentation is not considered discardable whitespace in Python.

Discuss limitations in Python due to Indent/Dedent tokens


2.1.3 Improved Representation of Magnets

The old format for describing magnets is difficult to read and worse to try to manually type. Some of its quirks are that magnets are separated by an unusual separator (`.:|:.:`), areas that accept other magnets (aka drop zones) are indicated by a seemingly random set of HTML tags (`
<!-- panel -->
`), and any special characters used in the code magnet must be escaped for HTML (so something like `1 < 2` would need to be entered as `1 < 2`). This format also is limited when trying to create more advanced types of magnets. The WAGS system addresses this problem by adding more form fields when adding “advanced Java magnets”, but this does not work well for adding additional languages.



This project creates an underlying object based data structure for magnets, as shown in figure . This structure can then be serialized as desired. Currently, serialization to the old magnet format is supported for backward compatibility, as well as JSON and YAML for more robust usages.

link figure of class diagram



2.2 What It Does: External Functions

2.2.1 Command Line Tool

This project provides a command line tool, which is a thin wrapper exposing the internal API functions to the command line.

2.2.2 Automated Interaction With WAGS Website

This project contains a proof-of-concept automated interaction with the WAGS website, which is capable of logging in to the site, and creating an assignment with magnets from a specified input file. The assignment that is created has the magnets loaded, but it does not have any of the other files that the WAGS system needs to automatically provide feedback to the students. The automated interaction is provided using libraries commonly used in Ruby for automated testing of websites. It uses the Capybara acceptance test framework (which provides a convenient, cross web driver DSL for interacting with websites), the Poltergeist web driver, and the PhantomJS headless browser. Support for Capybara, the Selenium web driver, and Firefox browser is coded, but Selenium will not trigger the log in button on the current web site.

2.3 What It Could Do: Future Extensions

2.3.1 Automatic Generation of Alternative/Distractor Magnets

The next step in easily creating code magnet assignments is to have the magnet creation tool not simply create magnets needed for the correct solution and any additional magnets specified by the instructor, but also to automatically create appropriate alternative “distractor” magnets for common student misconceptions and errors. This level of manipulation is available because we have the whole parse tree to work with during magnet creation, but would require significant work per language to define.

2.3.2 GUI Tool

A GUI tool that can open an input file, perform the actions of the CLI tool on it, and also has an editor to assist in the addition of any instructor directives

desired would be a nice addition to this project.

2.3.3 Further Automated Interaction With WAGS Website

It would be nice if this project could completely and reliably upload new magnet assignments to the WAGS website. The current interaction with the website does not upload the test file or any information other than the magnets. It also does not support the advanced Java magnet upload page. It also only works with the headless web driver (Poltergeist), not the one that shows the user what is happening (Selenium). However the limiting factor is that the functionality of the WAGS site is split between two versions of the website, and that the site has not been designed with programmatic interaction in mind.

2.3.4 Supporting New Languages

This project currently can create magnets for Python3 and Java, however it is set up to easily allow the addition of new languages. Basically, this is done by finding and adding an ANTLR4 grammar file for the desired language, making sure it conforms to a handful of guidelines, writing a short configuration file, and running a setup action on the project. Full details of this process are in section 4.5 [The Expansion: Adding a New Language](#).

2.3.5 Additional Serialization Formats

Because magnets are now objects, additional serialization formats (such as XML) could easily be added.

3 User Guide

3.1 Quick Start - CLI Tool

The primary interface to this project is the magnetizer command line tool. This tool is packaged as a stand-alone executable JAR file.

It takes as input a file name, and outputs the magnets created from the input file. By default it assumes the input file is in Java. If it is in Python, the language options needs to be specified as Python3.

By default, it prints the magnets in WAGS format to the console. It can also or instead print the magnets to JSON or YAML. It also can send the magnets to a specified output file, rather than the console.

```
Usage: magnetizer [options] file
  -l, --language LANGUAGE      Specify a language (default: Java,
                                other: Python3)
  --json                       Print the JSON output
  --yaml                       Print the YAML output
  --[no-]wags                 Print the output as WAGS magnets
```

(default)

```
-o, --output-file BASE_FILENAME  Output to a file
-h, --help                        Show this message
```

3.2 Installation for Use

This tool is packaged as a stand-alone JAR file. It does not need to be installed before running. Its only dependency is that Java needs to be installed, and it has been tested with Java ??

test Java
versions

3.3 Explanation of Magnet Creation

3.3.1 Java

The current configuration for Java creates magnets for

- Package declarations
- Import declarations
- Type declarations
- Class body declarations - this is the nonterminal that includes anything that can be directly in the class body
- Block statements - this is the nonterminal that includes anything that can be directly inside a block.

explain these
in a way
that doesn't
require deep
understand-
ing of the
ANTLR
grammar.

3.3.2 Python 3

The current configuration for Python creates magnets for

- simple statements
- compound statements

3.4 Modifying the Magnet Creation: Instructor Directives

Instructor directives are information that can be put in the input file that change how the magnets are created. They are a special version of comments. The triggering syntax varies per language, but the directives are the same for all languages.

Java `/*# <directive> */`

Python3 `## <directive>`

Note that directives in python are somewhat limited because comments cannot occur in a line.

Verify that
there are
no inline
comments in
python

explain
which direc-
tives work in
python

3.4.1 ALT and ENDALT: Creating Alternative Magnets

The ALT and ENDALT directives are used to create alternatives with different text for a magnet. The ALT directive also takes the desired alternative text, and the ENDALT directive is used to indicate where the end of the alternative text is. These directives should always be used as a pair, and encompass the text that should be replaced. They cannot go around anything that would trigger the creation of a new magnet or drop zone.

Concrete example of ALT

3.4.2 NODROP: Suppressing Drop Zones

The NODROP directive suppresses the creation of drop zones for the following magnet. An example of when this would be used is to put an entire loop on a single magnet, rather than having the individual statements in the body be on their own magnets.

Concrete example of NODROP

3.4.3 EXTRAMAG: Creating an Extra Magnet

The EXTRAMAG directive is used to add a magnet that is unrelated to any of the magnets in the input file.

Concrete example of EXTRA-MAG

4 Developer Notes

Targeted to someone expanding on this project

4.1 The Environment: Languages and Libraries Used

This project is written in Ruby for the JRuby interpreter. JRuby is able to access Java classes from Ruby. For this project, this means we can use to use a parser generated by ANTLR from Ruby.

Section needs more details

Automated interaction with the WAGS site is provided by using Capybara to interact with Selenium (drives Firefox) or Poltergeist/PhantomJS (headless).

Rake (Ruby make) is used to process new grammar files and other build tool functionality.

RSpec is used to handle testing.

4.1.1 Installation for Development

Consider if this section should be on its own

4.2 The Design

4.2.1 Directory Structure

- /bin** Contains the Ruby script that has set up as an executable and is the command line interface.
- /data** Contains grammar files and language-specific configuration information. This is the only location that has language-specific information.
- /doc** Contains documentation for the project.
- /etc** Contains any external projects and libraries needed for this project that are not available as Ruby gems.
- /java** Contains Java code generated by ANTLR and the compiled class files from that generated Java code.
- /lib** The primary directory of the project. This contains the Ruby code for the magnetizer.
- /spec** The second of the two folders used for testing. This is for RSpec spec tests.

4.2.2 JRuby / Java Integration

4.2.3 Key Classes

MagnetEmitterBase A mix-in that contains methods that have the same implementation across all the ruby-generated Visitors. In Ruby, a mix-in is a module that provides similar functionality to an abstract class in Java. When a class “include”s a module as a mix-in, that class gets all the methods defined in that module. Unlike abstract classes in Java, a class can have multiple modules mixed in, and a mix-in is unable to define abstract methods (Ruby does not have an equivalent to abstract methods).

MagnetEmitterVisitorGenerator Creates the correct subclasses of the ANTLR BaseVisitor. This class is where the bridging magic of JRuby is the most obvious, because we are creating subclasses of a Java class in Ruby.

ANTLR provides two primary patterns for traversing a parser tree: listener and visitor. This project uses the visitor pattern so that we do not have to visit child nodes in the parse tree for nodes that we know are terminal to our magnet generation.

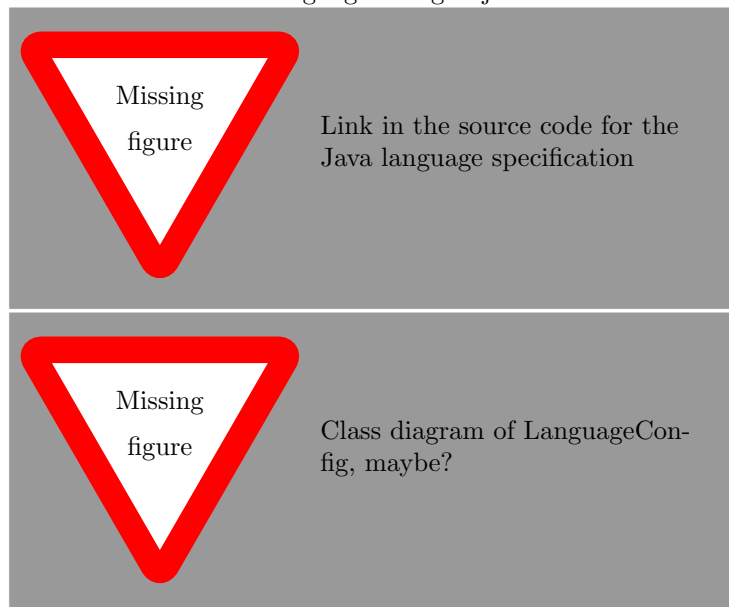
4.3 The Testing

4.4 The Configuration: Specifying Magnet Sections

The configuration that handles specifying how the parse tree for a language should be converted to magnets. The configuration is loaded from a YAML

file in the `/data` directory. A language configuration contains the name (which corresponds to the name of the grammar file), the name of the rule that is the start rule for the parsing, whether or not the language should have a file-level drop zone, information about the special comment that indicates directives for that language. The names of the rules that indicate are listed by which WAGS section the magnets will need to go into. There is also a section to specify overrides for cases where the name of the rule is not sufficient information, and there needs to be additional information used to determine which list a magnet should go in.

This YAML file describes instances of `LanguageConfig` class, and the magnetizer then uses the `LanguageConfig` objects to create the magnets.



4.5 The Expansion: Adding a New Language

To add a new language to be magnetized, a new grammar needs to be added, the project rebuilt to include that grammar, and a new configuration as mentioned in section 4.4 **The Configuration: Specifying Magnet Sections** needs to be added. The build file (Rakefile) also provides some additional tools that are helpful to visualize the new grammar when writing the configuration.

4.5.1 New Grammar Specification

A new G4 file can be added to the system. However, whitespace is required to go on channel 1, or your magnets will not have any whitespace, and you might get things like `"publicclassMyClass"`. Also, special comments for the directives need to be defined, and they need to go on channel 2. The entire comment for a directive should be recognized as a single token.

4.5.2 Rakefile: Useful Actions for Adding a New Language

References

- [1] Antlr Homepage.
- [2] antlr/grammars-v4.
- [3] Bison Homepage.
- [4] Creating Code Magnet Lab.
- [5] Cup Homepage.
- [6] YACC Homepage.
- [7] Terence Parr. *The Definitive ANTLR 4 Reference*. The Pragmatic Programmers, 2012.

Final export
of bibliog-
raphy from
Mendeley
and link it

Make sure
citations and
bibliography
are styled
correctly